

SYLTRA HEALTH — Regulatory, Data Governance & Privacy Document

Regulatory posture, consent architecture, and data ethics

Version 0.1, working draft. SYLTRA ONE. For advisors, investors and regulatory review.

o. What this document says

Here we explain how SYLTRA HEALTH complies, both regulatorily and ethically. Where we stand on medical-device classification, how we govern your consent and your data, and what we do not build before we meet the regulatory requirements.

Above all, what follows is a design posture and a working roadmap, not a legal opinion. We review every obligation with qualified regulatory counsel before any launch, and before any expansion of data scope or product line.

1. Where we stand today, and where we do not go early

We separate our product lines clearly, and each line has a gate we do not cross.

The line that is live today is the Wellness Safety line. It is wellness technology that helps with understanding and follow-up; we do not offer it as a medical device, and it does not diagnose or treat. It covers general integrations, understanding movement and environment, reminders and verification, the trusted circle, and medication adherence as a technical service.

The Connected Care line is deferred. Any feature that relies on compliant medical devices, clinician-defined plans, or deeper analytics is something we neither build nor speak of in the present tense before a documented review of Software-as-a-Medical-Device status and its clinical evaluation. The same is true of the Emergency Response line; any formal integration with response entities is something we do not build before regulatory requirements, agreements and testing are complete. The rule we hold to is that we describe no later line in marketing in the present tense before it passes its gate.

2. What we say the product does, and what it does not

We say it plainly, even in the interface itself. The platform detects unusual changes in the pattern, helps verify before escalation, alerts trusted people under your plan, and prepares the home environment with your consent. In turn, it does not diagnose or label a condition, does not predict crises, does not monitor medically around the clock, does not change a dose or treatment plan, and does not call emergency services automatically. These limits appear in the app and in every marketing asset, and we review them before any launch so that what we say stays within the wellness scope.

3. Consent as an architecture, not a checkbox

With us, your consent is an entity in its own right, not a box ticked once and forgotten. It is granted for each type of data on its own, for health, home, location, medication, and sharing with a specific person, and each consent has a scope, a grant date and a revocation date. When you withdraw it, the withdrawal takes effect at once, and we show you what stopped as a result. No one joins your trusted circle without accepting the role themselves, so no one is enrolled without their knowledge. And we ask for consent in advance and for a specific purpose; we do not gather data merely because it is available.

4. We collect the least necessary, and delete on request

We do not store a health field that a live feature does not use, and adding any new health field is something we state and justify. Where we can, we prefer to keep derived signals and events rather than a continuous raw stream. When you ask us to delete, we delete even the derived data, not only the raw. And every type of data has a defined retention period that we review with regulatory counsel.

5. A log you can see, and access that never bypasses your consent

We keep a sharing log always available to you, showing what was shared, with whom, when, and on what justification. Behind the scenes we keep an audit log for every access to sensitive data, recording the actor, the user, the action, the resource, and the associated consent. And every read or share path passes first through an active-consent check, so no query bypasses it.

6. How we protect the data

We encrypt data in transit and at rest, and manage secrets outside the code. We give each role the least privilege it truly needs. We separate development and test environments from production, and in development and testing we use only synthetic data, never real household data. And we align with the recognised information-security frameworks listed in the references.

7. Privacy is part of the design

We never share your individual data with any entity that prices a service or coverage, in any product line, whatever the return, and we build no endpoint that does so even if asked. When we share at a time of need, we hold to the least necessary, that something warrants reassurance without your health details, unless you have authorised a specific person. For us, privacy is part of the experience's design, not a legal paragraph at the bottom of the page.

8. People of determination and access, a right not a privilege

We build access into the heart of the product rather than as an add-on, and we treat the data of people of determination with the same care. Any setting that touches the access mode, sharing or contact requires explicit consent and a clear

record, and the governing principle is to grant greater independence, not greater surveillance. We never make access options a pretext for collecting extra data; there too we hold to the minimum, like everywhere else.

9. AI governance and accountability

Everything shown to the user, we can explain in one sentence, and any alert we cannot explain, we do not send, and we store the reason for each alert rather than the alert alone. Any model output stays advisory until it passes the safety rules, so the system takes no critical action on its own. Critical decisions and escalation stay within a plan that you control, not automated emergency decisions. And we continually evaluate data quality and any bias before any expansion.

10. What we complete before each gate

Before we move to any later line, we complete and document the Software-as-a-Medical-Device and risk classification with a regulatory review, carry out a data-protection impact assessment upon any expansion of data scope, sign data-processing agreements with any third party such as a pharmacy or messaging provider, establish documented validation and safety testing with a failure criterion written in advance, and run a textual review of every claim that appears in the interface or in marketing.

11. References and standards

The standards and laws below are real. Confirm the latest edition and scope of applicability with regulatory counsel before any reliance or publication.

1. Kingdom of Saudi Arabia. *Personal Data Protection Law (PDPL)* and its implementing regulations.
2. Saudi Data and AI Authority (SDAIA). *Personal data protection controls and guidance*.
3. Saudi Food and Drug Authority (SFDA). *Medical Devices Regulations (MDS)*.
4. IMDRF SaMD Working Group. *Software as a Medical Device (SaMD): Key Definitions*. 2013.
5. IMDRF SaMD Working Group. *Software as a Medical Device (SaMD): Clinical Evaluation*. 2017.
6. U.S. Food and Drug Administration. *Clinical Decision Support Software — Guidance for Industry and FDA Staff*. 2022.
7. ISO/IEC 27001. *Information security management systems — Requirements*.
8. IEC 62304. *Medical device software — Software life cycle processes*.
9. ISO 14971. *Medical devices — Application of risk management to medical devices*.
10. ISO 13485. *Medical devices — Quality management systems — Requirements for regulatory purposes*.

12. What remains open

A few things remain for us to settle with counsel. Account ownership when a caregiver manages it, where our starting position is that the patient owns and the caregiver is delegated a revocable scope. The minimum we actually store

versus what we process and then discard, per type of data. The requirements for keeping data within the Kingdom. And the scope of data-processing agreements with partners before the first operational integration.

End of document. It is a design posture and a working roadmap, not a legal opinion, which we review with regulatory counsel before every launch.